

Martian surface characterization using supervised machine learning

Utkarsh Mali,^{1,2}*

¹Department of Physics, University of Toronto, 60 St. George Street, Toronto ON M5S 1C6, Canada

²Canadian Institute of Theoretical Astrophysics, University of Toronto, 27 King's College Cir, Toronto M5S 3H8, Canada

Accepted XXX. Received YYY; in original form ZZZ

ABSTRACT

Sublimation in the polar ice regions occurs due to a buildup of CO₂ under the thick ice in the winter. In the spring, solar radiation causes the ice to rupture, sending the eroded dust airborne. The dust is blown by surface winds as it settles down causing detectable topographic features. I explore the use of supervised machine learning methods in detecting and characterizing these surface features. I begin by preprocessing the data using principal component analysis along with other preprocessing methods. I then implement a suite of supervised machine learning methods using the k-fold cross validation technique. In general, the accuracy scores were between 60% to 70%. The best performing model was a Gaussian naive Bayes with an accuracy of 71.8%. Upon further analysis we find the model to be skewing towards falsely predicting positive features. Over 20% of the items that were not features or "fans" were predicted as a feature. I conclude by providing potential followup studies. Overall, this project aims to highlight the potential for machine learning to aid the study of Martian surface recognition and its uses. By involving citizen scientists in the classification of these features, we aim to further engage the people in the study of Mars. In doing so, foster a greater understanding of the planet among the general public.

Key words: methods: statistical – planets and satellites: surfaces – planets and satellites: atmospheres

1 INTRODUCTION

The planet Mars has long been a subject of fascination for the both the public, and the scientific community. Studying its properties has peaked the interest of many great researchers. Modern instruments have greatly enhanced our ability to study the planet and its various features. One area of particular interest is the study of cold jets in the southern region of Mars, which has been shown to have an impact on the planet's atmosphere Kieffer (2007). These cold jets arise from the sublimation of carbon dioxide in the polar ice cap. The ice on the surface of the Martian surface is very thick, often over a meter in thickness in certain areas Kieffer et al. (2006). During the Martian winter, the carbon dioxide forms under the thick surface ice. As a result, high-pressure gas builds up beneath the slab. After winter, the temperature rises. The solar radiation impacting the gas under the ice builds up, eventually causing ruptures in the thick ice. Eventually, sublimation occurs. The ice ruptures and triggers a disruptive jet spraying the sub-ice carbon dioxide up into the air as a strong wind. These winds have been shown to be affected by the atmosphere. Kaufmann & Hagermann (2017). As the CO₂ jet releases, the pressure erodes the Martian surface dust and cases it to be blown up into the surface along with the CO₂. The air-borne dust slowly settles onto the Martian surface, it is affected by the surface currents and winds as it settles Aye et al. (2019). The imprint that is left on the surface carries information about the weather pattern at the time at which the sublimation occurred. These surfaces can be used to analyse the local seasonal wind which provides insight into

the Martian weather patterns.

In recent years, there has been an increase in the number of satellites and surveys observing Mars Sharma et al. (2021); Fisher et al. (2005). This greatly accelerated the study of the planet, and its various features. This in turn, has lead to an increase in the amount of data that needs to be processed by scientists Pan et al. (2017). The surveys generate vast amounts of data about the surface and atmosphere. This presents both challenges and opportunities. The sheer volume of data can be difficult to manage and analyse. On the other hand, it provides a wealth of information which can be used to train and infer new insights from. One way to address the challenge of managing large amounts of data, and their classification of features is to involve citizen scientists Bird et al. (2018); Banerji et al. (2010). Doing this speeds up the process of data analysis and makes it more efficient. Citizen scientist projects like WISE and SpArcFiRe have already demonstrated the potential for involving the public in academic research Peng et al. (2018); Nguyen et al. (2018). The Planet Four Collaboration is a citizen science project. First initiated in 2014 by a team of researchers at the University of Arizona. It focuses on the topographic features of the Martian surface, such as fans, streaks and blotches that appear during the sublimation process. By recruiting volunteers to classify these features, the collaboration was able to generate large amount of data which can be used to train machine learning models. In doing this, the collaboration will be able to greatly speed up the process of data analysis. In addition, off loading the predictive task to a machine learning model will enable scientist to focus on the identifying macroscopic patterns in the data. This is more likely to provide

* E-mail: utkarsh.mali@utoronto.ca

valuable insight about new science.

In my exploratory study, I will explore the extent to which methods in supervised machine learning (ML) can aid the classification of Martian surface features. To accomplish this, I begin by formatting the images into a standard format. I then clean and preprocess the image data to enable better training of the ML models. This includes data augmentation, gamma correction, principal component analysis (PCA) and histogram oriented gradients (HOG). Once cleaned, I apply multiple supervised ML models to train the data. Some of the models are frequentist; logistic regression (LR), support vector machines (SVM), k-nearest neighbours (KNN), while others are Bayesian; linear/quadratic discriminant analysis (LDA/QDA), naive Bayes (NB), Gaussian process (GP). Finally, I explore the possibility of applying unsupervised methods such as stochastic gradient descent (SGD). Once trained I evaluate the model performance with accuracy scores and study the best model in depth through a confusion matrix.

2 DATA

The data for this project was measured on the Mars Reconnaissance Orbiter (MRO). It was a spacecraft launched by NASA in 2005 with the mission of mapping important regions in Mars. On it contained many cameras and feature detection devices, one of which are the *High Resolution Imaging Science Experiment* or HiRISE for short. It is used to study geology and surface features on Mars with a resolution of up to 0.3 meters per pixel McEwen et al. (2007); Simpson et al. (2014). Over its commission the HiRISE instrument focused on a few regions of interest in the south polar region, these images captures the surface features highlighted in Section 1. The photos taken are open source and can be found on the Zooniverse [Planet Four Collaboration](#). The catalogue includes images of the Martian surface along with the classifications of features marked by citizen scientists. The numerical feature information can be converted into boxes represented in a sample tile Fig 1. Each tile is uniquely identifiable with many classifications per tile, representing multiple fans per image. In supervised machine learning, the data is typically split into training and testing data. The training data is used to generate reference samples for a model while the testing data is used to evaluate the models performance. Doing this enables a model to perform well on new, unseen data. For each tile, random boxes are generated within the image, these random boxes, both random in size and position are known as "notfans". The user-marked features are labeled "fans". Together, they create ground truth and false categories. In our model, we aim to predict if the feature is a fan given a random box. Our train/test split, highlighted in Fig 2, consists of 1400 train and 1400 test photos. We skew our dataset towards "notfan" classifications with the aim of eliminating false positives. With the data evenly split between random boxes and user classified fans, we preprocessing the data in order to denoise and improve train time.

3 PREPROCESSING

Preprocessing images involves using a set of methods to convert the image into a standardized format. It also includes feature extraction, which identifies important characteristics from the images that the machine learning model can use. The goal of this process is to improve the performance of the model. In this section, we will outline the steps we take to achieve this goal.

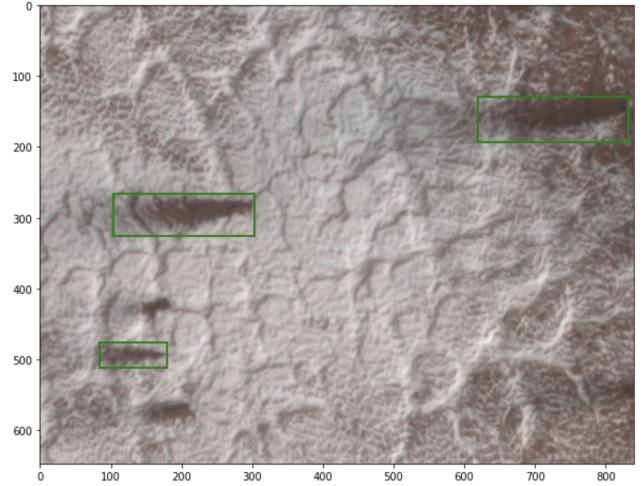


Figure 1. Sample image of a single "tile" in the data taken from a larger image with multiple tiles. The image highlights topographical features on the Martian surface. Each green box represents a marked item on the surface with a specified classification id.

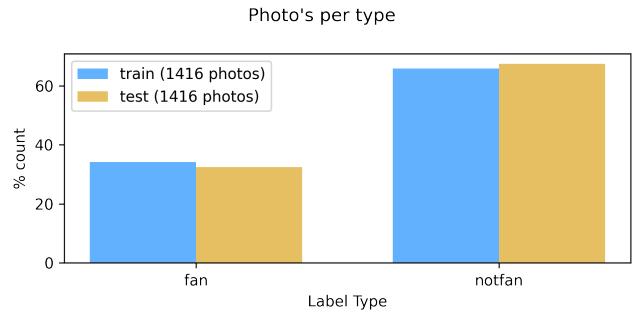


Figure 2. Representation split of the dataset into training and testing data. The size of the data was reduced due to limitations in training time.

3.1 Image Augmentation and Gamma Correction

We begin our analysis pipeline with data augmentation. We use this to increase the size and diversity of the training data. As our dataset is limited, data augmentation may help prevent over-fitting and improve the general performance of the model. We do this by translating, scaling and flipping "fans" to create a larger data space with new augmented versions of the fan images. An sample of the data augmentation pipeline is highlighted in Fig 3 in which a same image is flipped and translated. In order to further improve the preprocessing pipeline we apply gamma correction to the images used in the training set. Doing this allows us to improve the brightness and contrast of the images. The equation is given by the following:

$$I_c = I_o^{\frac{1}{\gamma}} \quad (1)$$

Here, I_c is the corrected fan image, I_o is the original fan image. This non-linear equation maps the pixel values in the original image to new values in the corrected image according to a power function, with the gamma value determining the exponent of the function. The gamma value controls the amount of correction applied, with higher values resulting in brighter fans and lower values resulting in darker

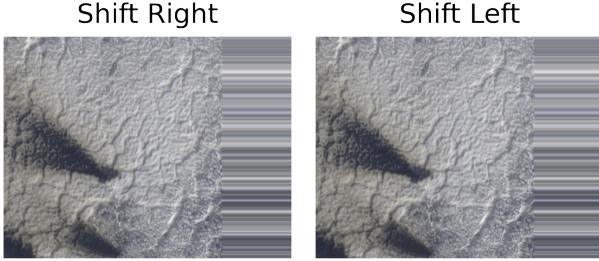


Figure 3. An example of a typical data augmentation process. In this case, both images are taken from a single reference image. The outputs contains multiple images with translated, flipped and inverted version of the reference image.

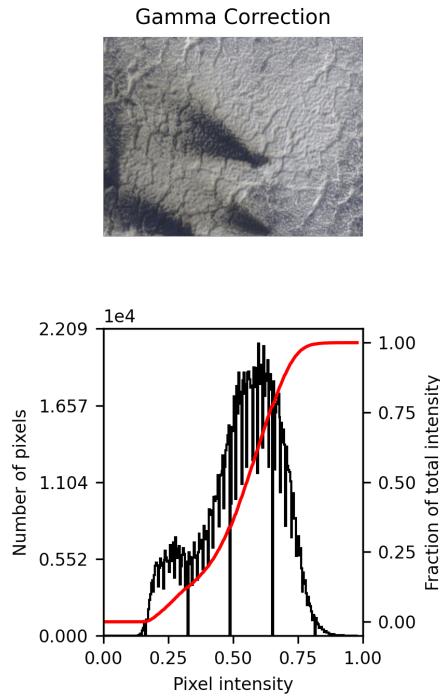


Figure 4. Image correction applied to the reference image in order to simplify further preprocessing. The x-axis represents pixel intensity and the y-axis represents count data. The red line represents the corrective fix applied to the image.

fans, in our case $\gamma = 2$. An example of this process is demonstrated in Fig 4.

3.2 Principal Component Analysis

After preprocessing the data, we reduce its size to improve the efficiency of our machine learning model. Initially, we considered using general data reduction techniques such as principal component analysis (PCA) and histogram oriented gradients (HOG). PCA is a mathematical technique in linear algebra that offers versatility and flexibility, while HOG is a feature extraction method. It divides an image into cells that generate gradients of orientations for each cell. These gradients are combined into a feature vector which represents

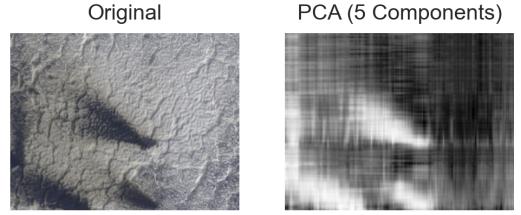


Figure 5. Original (top) vs reconstructed (bottom) image with the application of principal component analysis (PCA). The reconstructed image still contains most of the features surrounding areas of interest while greatly reducing the size of the data.

the entire image. We ultimately choose to use PCA for its ability to denoise images and its superior predictive performance. We will discuss this technique in more detail below.

When applying PCA we use the single vector decomposition (SVD) method, a mathematical method to reduce the original matrix into a set of diagonal and unitary matrices. We then make an orthogonal projection of the components onto a diagonal matrix order by the variance. The diagonal elements of the decomposition are known as the principal components. We apply PCA on the x-y image domain. [Mudrova & Procházka \(2005\)](#)

$$\mathbf{I}(x, y) = \mathbf{V}\boldsymbol{\Sigma}\mathbf{U}^\top \quad (2)$$

Here \mathbf{I} is the image. \mathbf{V} is a unitary matrix and is known as the *left singular matrix*. We use \mathbf{V} to project vectors from the image to feature (component) basis. $\boldsymbol{\Sigma}$ is a diagonal matrix ordered by decreasing magnitude (i.e. $\sigma_1 \gg \sigma_2 \gg \dots \gg \sigma_n$). Here σ_i is known as a singular value which holds information about the data vector. \mathbf{U} is also a unitary matrix. It is known as the *right singular matrix*. Using this we form the feature vector. We do this by projecting the image vector \mathbf{i}_i onto the change of basis vector \mathbf{V} to get the principle components of the input data. The feature vector is also known as the principal components of the transformation. An example of the reconstructed image is shown in Fig 5. The reconstructed images reduce noise and make the fan images easier to train on.

4 MODEL SELECTION

When selecting a supervised ML model, multiple factors, like size, and data quality must be taken into consideration. We study the significant impact that model choice has through k-fold cross validation. This occurs by diving the dataset into k subsets (folds). The model is then trained ($k = 10$) times using a unique data subset as a validation set while training on the remaining subsets as training sets. Comparisons between each model use multi-threading to speed up compute time. The final performance metrics are calculated using an average of the model performance over each subset. We use this to tune hyperparameters for each model and aid our decision in model selection. In model selection, we restrict our search to mainly supervised learning methods. [Cady \(2017\)](#)

4.1 Model: Support Vector Machines

A support vector machine (SVM) is a classification algorithm that splits the training data over a decision boundary. This boundary maximally separates the two classes, fans and not fans. It is represented mathematically in the following form:

$$\min_{w,b} \frac{1}{2} w^T w + C \sum_{i=1}^n \max \left(0, 1 - y_i (w^T \phi(x_i) + b) \right) \quad (3)$$

The decision boundary is represented as a hyperplane of PCA components with w being the weight vector normal to the hyperplane. Here, b is the bias term and C is the penalty intensity parameter. C controls the trade-off between a large classification penalty and no classification penalty. The former represents data over-fitting while the latter represents under-fitting. The equation is solved as a quadratic optimization problem with linear constraints. It can be solved using the Lagrange multiplier method. Once the decision model has been trained, it can make new predictions by computing the value of the decision function $f(x) = w^T x + b$ for each example. Doing this allows the model to assign a new image the appropriate fan/notfan class. Since we are solving a binary classification problem. Images with $f(x) > 0$ are classified as one class while images with $f(x) < 0$ are classified as the opposite class. This algorithm has a run-time complexity of $O(n^3)$ although lower values of C may reduce this [Hearst et al. \(1998\)](#).

4.2 Model: k-Nearest-Neighbours

k-Nearest-Neighbours (kNN) is a simple classification algorithm. New images are predicted based on the classification of k nearby labeled images. The algorithm finds k examples in the training set which are the nearest (have the closest PCA vector) to the new example. When computing the distance, we used an L2 norm. Formally:

$$y^* = \max_{t^{(z)} \in \text{class labels}} \sum_{i=1}^k \mathbb{I}(t^{(z)} = t^{(i)}) \quad (4)$$

In which we have found k examples $\{\mathbf{x}^{(i)}, t^{(i)}\}$ nearest to the test instance \mathbf{x} . The kNN algorithm is simple and easy to implement, it requires no training phase (i.e. is used directly on the dataset). However, kNN runs into severe computational complexity problems. The requirement in which it must compute the L2 norm between a new example and all examples in the training set makes it struggle with high-dimensional data. The distance between each data-point approaches infinity as the number of PCA components (dataset size) is increased [CSC311 \(2022\)](#).

4.3 Model: Logistic Regression

Logistic regression (LR) is an algorithm which is typically used for binary classification tasks. The model learns a set of parameters that define the logistic function. A function which splits the fan/notfan classes in binary classification. Once trained the function maps the input space to the output space using a monotonically increasing function. The logistic function is defined as the following.

$$f(x) = \frac{1}{1 + e^{-k(x-x_0)}} \quad (5)$$

Here, x_0 is the midpoint value and k is the logistic growth rate. The logistic function outputs a probability between 0 and 1 that the test image belongs to a certain class. We set the classification threshold

to 0.5. The model is optimized by implementing a penalty between probability of a certain prediction and the ground truth value of that image. This is written formally below:

$$\begin{cases} -\ln p_k & \text{if } y_k = 1 \\ -\ln(1-p_k) & \text{if } y_k = 0 \end{cases} \quad (6)$$

This equation can be turned into a loss function by combining the terms to obtain the *cross-entropy*, and then maximizing the inverse log-likelihood. The benefit in using LR is that the threshold may be changed above/below 0.5 in order to eliminate false positives/negatives [CSC311 \(2022\)](#).

4.4 Model: Naïve Bayes

Naïve Bayes (NB) is a probabilistic classifier that applies Bayes theorem with a strong assumption that the feature vectors are independent. Only a small amount of data is required for parameter classification. The conditional independence assumption allows us to rewrite a version of Bayes theorem.

$$\hat{y} = \arg \max_y P(y) \prod_{i=1}^n P(x_i | y) \quad (7)$$

Here, y , is the class fan/notfan, and x is the PCA vector space. The probability $P(y)$ is given by the training ground truth.

$$P(x_i | y) = \frac{1}{\sqrt{2\pi\sigma_y^2}} \exp\left(-\frac{(x_i - \mu_y)^2}{2\sigma_y^2}\right) \quad (8)$$

Our version implements Gaussian Naïve Bayes which uses a Gaussian likelihood $P(x_i | y)$ in its approach. Here σ_y and μ_y are used in estimating the maximum likelihood. In this approach, we exploit the fact that NB is a quick yet simple classifier. [Kaur & Oberai \(2014\)](#)

4.5 Model: Gaussian Processes

A Gaussian process (GP) is a stochastic method in statistics used for non-parametric regression and classification. Since it does not make any assumptions about the underlying data distribution, it is robust when trained on smaller data sets. It is able to model the relationship between the input, PCA components and output, classification fan/notfan using a Gaussian distribution. This continuous distribution is defined by its mean and covariance. The GP model uses a mean function to model the conditional mean of the fan/notfan classification given the PCA components. In random variable notation, it has the following form:

$$y(x) \sim GP(X, K(X, X')) \quad (9)$$

Here X is the mean function and $K(X, X')$ is the covariance kernel. The key predictive equations are shown below.

$$\bar{f}_* \triangleq \mathbb{E}[\mathbf{f}_* | X, \mathbf{y}, X_*] = K(X_*, X) [K(X, X) + \sigma_n^2 I]^{-1} \mathbf{y} \quad (10)$$

$$\text{cov}(\mathbf{f}_*) = K(X_*, X_*) - K(X_*, X) [K(X, X) + \sigma_n^2 I]^{-1} K(X, X_*) \quad (11)$$

The former equation represents the functional form of the mean maximum a posteriori (MAP) while the latter equation represents the optimal covariance. Together they are used to make new classifications on new datasets (PCA components). As with Logistical

Regression, we are able to tune the threshold probability to prevent false positives from arising. A benefit to GP classification is that we are able to propagate the errors through the training set. One drawback to GP classification is the $O(n^3)$ run time. Similar to support vector machines, the inversion of the covariance matrix results in a large computational cost added to it. Its popularity in astronomy has made it a viable candidate to investigate. Williams & Rasmussen (2006)

4.6 Model: Linear Discriminant Analysis (and Quadratic Discriminant Analysis)

Both Linear Discriminant Analysis (LDA) and Quadratic Discriminant Analysis (QDA) are classification algorithms which attempt to split the data into the fan/notfan classes. LDA does this by finding the best linear combination which maximize the separate between fan and notfan. This is done by optimizing the direction in which the data varies the most and projecting the data onto that direction. QDA is similar to LDA, but instead of linear boundaries between the fan/notfan class, quadratic boundaries are allowed. This enables QDA to find more complex patterns in the data. A side effect to this is that as a result of this, the data may also be over fitted. Formally, $p(\vec{x}|y = \text{fan}) = (\vec{\mu}_{\text{fan}}, \Sigma_{\text{fan}})$ and $p(\vec{x}|y = \text{notfan}) = (\vec{\mu}_{\text{notfan}}, \Sigma_{\text{notfan}})$. Below is the Bayes optimal solution to predict points given the conditional probability density functions are normally distributed.

$$\begin{aligned} & (\vec{x} - \vec{\mu}_{\text{fan}})^T \Sigma_{\text{fan}}^{-1} (\vec{x} - \vec{\mu}_{\text{fan}}) \\ & + \ln |\Sigma_{\text{fan}}| - (\vec{x} - \vec{\mu}_{\text{notfan}})^T \Sigma_{\text{notfan}}^{-1} \\ & (\vec{x} - \vec{\mu}_{\text{notfan}}) - \ln |\Sigma_{\text{notfan}}| > T \end{aligned} \quad (12)$$

This is known as the QDA. We can reduce this to LDA by assuming $\Sigma_{\text{notfan}} = \Sigma_{\text{fan}} = \Sigma$. While LDA is generally preferred in the ML community due to its efficiency and non-overfitting capabilities. Alternatively, QDA, due to its ability to choose priors can be generally considered a Bayesian classification method. We choose to apply both LDA and QDA. Our feature space is small (only a few PCA components are used) compared to the large number of data points, as a result, over-fitting should not occur.

4.7 Model: Stochastic Gradient Descent (Unsupervised)

As an extension to the supervised methods used, Stochastic Gradient Descent (SGD) is an unsupervised algorithm designed to find the minimum value of a function. It is an iterative algorithm that starts with an initial guess, and takes small steps in the direction of the steepest negative gradient. The size of these steps are known as the learning rate and the determination of what the value should be is called the loss function. The minimization of this loss function provides the classification of the fan/notfan. In our case, we use a hinge loss function.

$$\ell(y) = \max(0, 1 - t \cdot y) \quad (13)$$

In this equation, t represents the class similarity and is usually ± 1 , y represents the classifier score. We choose this due to its similarity to SVMs which are well understood in the literature. We expect the SGD to be a quicker version of the SVM classifier. While slightly outside the scope of this study, it is useful to study the SGDs performance in relation to that of the supervised methods listed above.

5 RESULTS

The results of the model training is presented in this section. We tested different models with varying hyperparameters for each proposed model. The models training was focused on proof of concept, understanding and pipeline flow over the direct results.

The first method used was Logistic regression, in this model we applied an L2 penalty with an saga optimizer. We were able to achieve a test accuracy of 69.4%. The model took 12 seconds to train. The version we applied was regularized logistic regression.

Next we applied Gaussian Process classification. In this classifier, we used an RBF covariance matrix Williams & Rasmussen (2006) with a length-scale of 1, the Laplace approximation was used to approximate a non-Gaussian posterior into a functional normal distribution. The test dataset resulted in an accuracy of 60.5%. The inversion of the covariance kernel resulted in a long training time of 168 seconds.

The next model used was Support Vector Machines. We applied the linear kernel version due to its convex quadratic optimization, we set the penalty term $C = 1$. The model performance resulted in an accuracy of 66.9% with a run time of 58 seconds.

Both Linear Discriminant Analysis and Quadratic Discriminant Analysis were applied with similar success, the former had a test accuracy of 62.2% while the latter had an accuracy of 65.8%. Both took approximately 6 seconds to complete training. LDA used singular value decomposition (SVD) to propagate the vectors into the diagonal space. There were no priors used in QDA, although this may be a possible extension to consider.

The KNN classifier was set-up using 5 neighbours with a uniform weight on each neighbour. The distance measured was Euclidean (L2 norm). The resulting test accuracy was 65.7% with a train time of 5 seconds. This model did return the highest variability in accuracy between k-folds of training.

Stochastic Gradient Descent was the only non-supervised approach used in training. As mentioned before, a hinge-loss model was used to imitate the performance of a SVM model. A maximum of 1000 iterations of gradient descent were set with a stopping criteria tolerance of 0.01 between steps. The model trained in less than a second with a predictive performance of 65.9%. Similar to KNN, this model had relatively high variability between the k-folds of training.

Naive Bayes was the final process used. Similar to QDA, it is a Bayesian classification method. Our model performed efficiently training is less than a second and providing the best accuracy of 71.8%. Gaussian Naive Bayes was used with small variational smoothing = $1e-9$.

The model was trained in python using mainly the `sklearn` package. Some dependencies occur in C and FORTRAN. The use of multi-threading played a large role in speeding up both the data wrangling and model training. This occurred through `joblib`. The plots were generated using `matplotlib` and `seaborn`.

A tabular representation of the results is shown in Table 1, with its corresponding figure comparison in Fig 8. A sample

Table 1. Test accuracy’s of each model performance averaged over k-folds ($k = 10$) of cross-validation. The best performing model was the Gaussian Naive Bayes (NB) with an accuracy of 72% and runtime of under 1s. The worse performing model, in both time and accuracy complexity, was the Gaussian Process Classifier (GP) with an accuracy of 61% and runtime of just under 3 minutes. All the remaining models has approximately similar performances. For reference, LR (Logistic Regression), SVM-I (Support Vector Machine - Linear), LDA/QDA (Linear/Quadratic Discriminant Analysis), KNN (K-Nearest Neighbours) and SGD (Stochastic Gradient Descent)

Model	Test Accuracy (out of 1)	Runtime (seconds)
LR	0.694	12s
GP	0.605	168s
SVM-I	0.669	58s
LDA	0.622	7s
QDA	0.658	6s
KNN	0.657	5s
SGD	0.659	1s
NB	0.718	1s

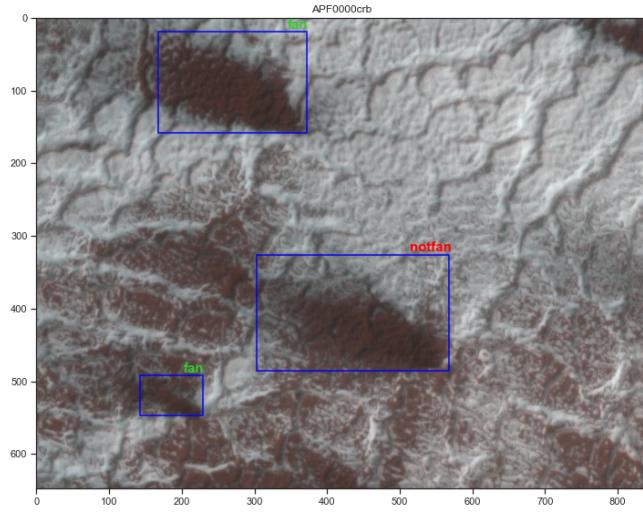


Figure 6. Sample tile with over-plotted results containing model predictions of the trained model. Green text corresponds to correctly predicted features, while red text corresponds to incorrectly predicted features.

image of the best Gaussian NB model performance can be seen in Fig 6. Finally, the confusion matrix representing the Gaussian NB performance for recall, precision and accuracy is shown in Fig 7. The model performs well at predicting the True Positive, and True Negative. However, it is also skewed towards over-predicting false positives. I.e. predicting that an image is a fan when a human has classified it as not a fan.

6 DISCUSSION

We aim to explore some of the key issues that arose in our analysis. We will begin the discussion with the challenges and limitations of our implemented models. We will then move onto some changes that we recommend. Finally, we will highlight possibilities for future work. We begin with the difference in performance of QDA vs LDA.

QDA is a generalization of the popular LDA classification al-

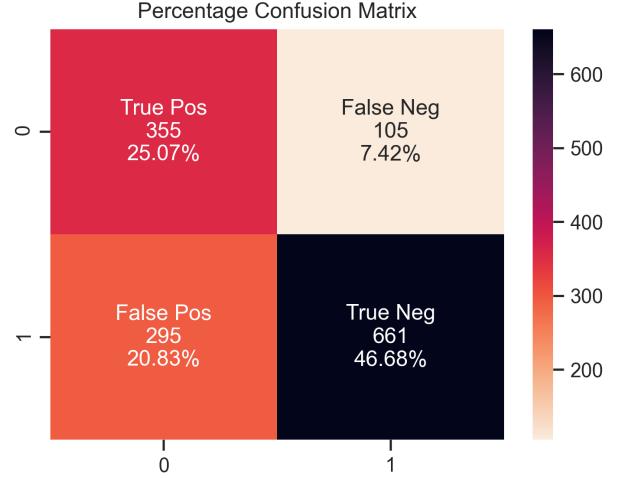


Figure 7. Confusion matrix showing a visual representation of the classification algorithm. The four entries of the matrix, in both percentage and count represent the True Positive (TP), False Negative (FN), False Positive (FP) and True Negative (TN) values. A False Positive corresponds to the model predicting a fan while the human classification is that the feature is not a fan. A False Negative corresponds to the model predicting that a feature is not a fan while the human classification is fan. While the model performs well in predicting true positives and true negatives, it has a skew towards predicting false positives.

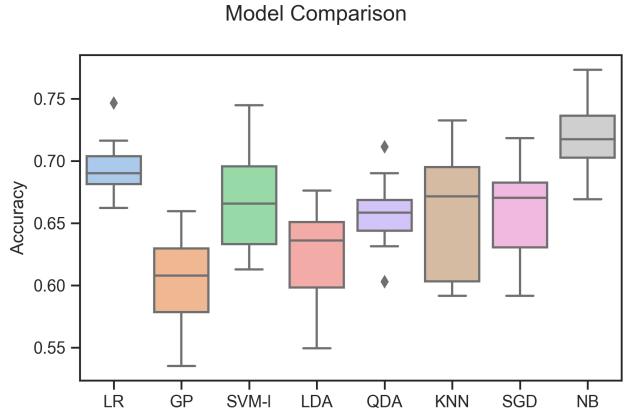


Figure 8. Pictorial representation of the tabular results highlighted in Table 1. As before, Naive Bayes was the best performing model with Gaussian Process being the worse performing model. While the rest of the models had similar performance, their spread greatly varies. Quadratic Discriminant Analysis and Logistic Regression both had small spread in their test accuracy’s. Models such as K-Nearest Neighbours and Support Vector Machines, had a larger spread in its accuracy. This will be further explored in the discussion.

gorithm. It was found in Fig 8 that QDA outperformed LDA, as expected. This was partly due to the nature of the data, but was also due to the nature of QDA. Since it models the covariance of each fan/notfan class separately it better captures the complexity of each data class, this results in more accurate predictions. On the other hand, LDA assumes identical covariance between each class resulting in a limitation which is not necessarily true. Trivially, random boxes will have different covariance to fans. Overall, the

results indicate that QDA is marginally more effective than LDA. This will further improve with the addition of class specific priors.

In our tests, we found in Table 1 that SGD was surprisingly outperformed by SVM. This was unexpected, especially since they used analogous loss-optimization functions. Upon further examination, we conclude that the use of strict learning rates, a small number of iterations and early stopping may have resulting in poor performance. Additionally, the PCA components used may have been sub-optimal for gradient descent learning. Overall, our results demonstrate the importance of carefully choosing the SGD hyperparameters as well as considering the training time and number of PCA components.

The choice of PCA components varies the complexity of the dataset being used. Increasing the number of components greatly changes the model performance. While giving the model more feature information to train on, it also includes more noise. After testing multiple models between 5-200 PCA components, we concluded that using as few as possible resulted in the best performance. In doing so, we were able to maintain most of the images key features while denoising the feature space. An example of the PCA is shown in Fig 5. Ultimately, the choice of PCA components is fluid and should be studied concurrently with model/hyper-parameter variations. PCA, however, is not the only method of dimensional reduction we explored. The HOG (Histogram Oriented Gradients) transform is widely used in computer vision when extracting features from fan images. However, when we attempted to use this transform in our pipeline for modeling training, we found that its performance was not a great improvement over PCA. This could be due to a variety of reasons, such as the quality of the input fan image, the nature of the fans themselves or the hyperparameters used in the HOG transform. Experimenting with this is a possible avenue of further study.

Another approach that was considered was using edge detection alongside a Fourier transform to extract the fan structure from the image. By doing this, we expect a peak in Fourier space around the edge of the fan structure. The difficulty in doing this would be to convert the pixel values of the edge to vector values. As a result, this preprocessing pipeline would probably need to be supplemented with a HOG transform stage before applying the Fourier transform. While this approx. does increase the preprocessing complexity, it could in theory provide a large performance improvement.

The results in the confusion matrix in Fig 7 clearly indicate a model skew towards classifying items that are not fans as fans. This is particularly harmful in certain situations. When trying to predict the Martian weather patterns, a false classification would result in an incorrect weather vector, deteriorating the weather predictive precision. It is possible to alleviate this issue by adjusting the some of the models threshold for predicting positive outcomes. For example, one may increase the threshold from 0.5 to 0.8 in both LR and GP. This would skew the data towards False Negatives which are much preferred over false positives. Additionally, using different NB classifiers may further improve the predictive performance.

The high False Positive rate demonstrates the need to increase the size of the fan dataset. While we have already applied basics such as cropping and flipping the images, simple methods like these have their limitations. One could consider using more advanced models such as generating new images using generative models or synthesizing a combination of images using computer vision.

Additionally, one may consider using different data augmentation techniques separately for the fan and notfan classes. As they are inherently have different data composition, they may benefit from separate treatment. Overall, the extension of different data augmentation methods should further improve the False Positive rate along with making the model more robust.

Commonly used in both industry and academia, one may consider using a convolution neural network (CNN) to classify fans instead of supervised methods. A potential advantage to this is that CNN's are able to extract and combine feature vectors in a hierarchical manner. This is particularly useful in image classification. Another potential advantage to this is that we are able to train the CNN's on increasing large amounts of data. Combined with better data augmentation this would greatly improve predictive performance and is an avenue that should be seriously considered.

This work, once streamlined, will be able to help predict Martian weather patterns. This in turn could inform decision making about missions to the planet. Furthermore, this work can be expanded to include other data sources such as rover observations, and spectroscopic samples. In doing so we may predict when and where the ice may rupture, and thus know where to point the satellites to measure data. Overall, this work will improve our understanding of the winter/spring ice sublimation and how to plan future missions around it.

7 CONCLUSION

The use of machine learning, both supervised and unsupervised has the potential to greatly improve our understanding of the Martian topographical features. By studying them we will be able to better understand the Martian surface. We implement a suite of machine learning models, attempting to detect and characterize these Martian surface features. In doing so, we implement a preprocessing pipeline that implements data augmentation, gamma correction and principal component analysis. We then train on multiple different machine learning models, mainly supervised (both Bayesian and frequentist) and compare their performance. The Gaussian Naive Bayes classifier produced the best accuracy and time efficiency metrics getting 71.8% of its predictions rights after training in under 1 second. Upon further analysis, we determined the main are of concern being the high false positive rate ($FP = 20.83\%$). We discuss possible ways of mitigating this, such as higher thresholds for our models and the implementation of mode robust preprocessing algorithms. We then highlight extensions to this work with unsupervised methods such as convolution neural networks. This project has demonstrated the feasibility of using supervised learning methods to identify Martian surface features. Further research in this area has the potential to yield valuable insights into the winter/spring ice patterns as well as aid future missions.

ACKNOWLEDGEMENTS

I acknowledge Professor Joshua Speagle for useful discussion and his valuable insight into statistical methods and Alexander Laroche for his discussion about Fourier transforms during my presentation. I acknowledge the NASA Mars Orbiter (MRO) and its HiRISE instrument with which the photos were taken. Finally I wish to acknowledge this land on which the University of Toronto operates. For thousands

of years it has been the traditional land of the Huron-Wendat, the Seneca, and the Mississaugas of the Credit. Today, this meeting place is still the home to many Indigenous people from across Turtle Island and we are grateful to have the opportunity to work on this land.

DATA AVAILABILITY

The data is part of the [Planet Four Collaboration](#) and is completely open source. The specific training images are available to download from the [results page](#) which contain both the fan catalogue and the observational metadata. More details about the files can be found [on this information link](#).

REFERENCES

- Aye K.-M., et al., 2019, Icarus, 319, 558
 Banerji M., et al., 2010, Monthly Notices of the Royal Astronomical Society, 406, 342
 Bird R., et al., 2018, Muon Hunter: a Zooniverse project, doi:10.48550/ARXIV.1802.08907, <https://arxiv.org/abs/1802.08907>
 CSC311 C., 2022, CSC 311 fall 2022: Introduction to Machine Learning, https://www.cs.toronto.edu/~rahulgk/courses/csc311_f22/index.html
 Cady F., 2017, The data science handbook. John Wiley & Sons
 Fisher J. A., et al., 2005, Journal of Geophysical Research: Planets, 110
 Hearst M. A., Dumais S. T., Osuna E., Platt J., Scholkopf B., 1998, IEEE Intelligent Systems and their applications, 13, 18
 Kaufmann E., Hagermann A., 2017, Icarus, 282, 118
 Kaur G., Oberai E. N., 2014, International Journal of Computer Science and Mobile Computing, 3, 864
 Kieffer H. H., 2007, Journal of Geophysical Research: Planets, 112
 Kieffer H. H., Christensen P. R., Titus T. N., 2006, Nature, 442, 793
 McEwen A. S., et al., 2007, Journal of Geophysical Research: Planets, 112
 Mudrova M., Procházka A., 2005, in Proceedings of the MATLAB technical computing conference, Prague.
 Nguyen T., Pankratius V., Eckman L., Seager S., 2018, Astronomy and computing, 23, 72
 Pan L., Ehlmann B. L., Carter J., Ernst C. M., 2017, Journal of Geophysical Research: Planets, 122, 1824
 Peng T., English J. E., Silva P., Davis D. R., Hayes W. B., 2018, Monthly Notices of the Royal Astronomical Society, 479, 5532
 Sharma M., Gupta A., Gupta S. K., Alsamhi S. H., Shvetsov A. V., 2021, Drones, 6, 4
 Simpson R., Page K. R., De Roure D., 2014, in Proceedings of the 23rd international conference on world wide web. pp 1049–1054
 Williams C. K., Rasmussen C. E., 2006, Gaussian processes for machine learning. Vol. 2, MIT press Cambridge, MA

This paper has been typeset from a $\text{\TeX}/\text{\LaTeX}$ file prepared by the author.